

B461 Database Concepts

Assignment 6

Fall 2023

This assignment tests the following concepts:

- Nested Relations and Semi-Structured Databases
- PL/pgSQL

To turn in your assignment, you will need to upload to Canvas a single file with name `assignment6.sql` which contains the necessary SQL statements that solve the problems in this assignment. The `assignment6.sql` file must be so that the AI's can run it in their PostgreSQL environment. You should use the `assignment6Script.sql` file to construct the `assignment6.sql` file. (note that the data to be used for this assignment is included in this file.) In addition, you will need to upload a separate `assignment6.txt` file that contains the results of running your queries.

1 Nested Relations and Semi-structured databases

Consider the lecture on Nested relational and semi-structured databases. In that lecture, we considered the `studentGrades` nested relation and the `jstudentGrades` semi-structured database, and we constructed these using a PostgreSQL query starting from the `Enroll` relation.

1. Write a PostgreSQL view `courseGrades` that creates the nested relation of type $(\text{cno}, \text{gradeInfo}\{(\text{grade}, \text{students}\{(\text{sid})\})\})$. This view should compute for each course, the grade information of the students enrolled in this course. In particular, for each course and for each grade, this relation stores in a set the students who obtained that grade in that course. **[15 points]**
2. Starting from the `courseGrades` view in Problem 1, solve the following query: Find each (s, C) pair where s is the sid of a student and C is the set of cnos of courses in which the student received an 'A' or a 'B' but not a 'C'. The type of your answer relation should be $(\text{sid} : \text{text}, \text{Courses} : \{(\text{cno} : \text{text})\})$. **[15 points]**
3. Write a PostgreSQL view `jcourseGrades` that creates a semi-structured relation which stores jsonb objects whose structure conforms with the structure of tuples as described for the `courseGrades` in Problem 8. Test your view. **[15 points]**

2 Object Relational Programming

The following problems require you to write object-relational programs. Many of these require programs written in Postgres' PL/pgSQL database programming language.

4. Write a PL/pgSQL function that takes an array of integers as input and returns a new array where each element is the sum of all the elements in the input array up to and including that element. For example, if the input array is $\{1, 2, 3, 4\}$, the output array should be $\{1, 3, 6, 10\}$. **[15 points]**
5. Write a function that takes a positive integer as input and returns an array of all the prime numbers up to and including that integer. **[15 points]**

6. Consider a parent-child relation $PC(\text{parent}, \text{child})$. (You can assume that PC is a rooted tree and the domain of the attributes parent and child is int.) An edge (p, c) in PC indicates that node p is a parent of node c . Now consider a pair of nodes (m, n) in PC (m and n may be the same nodes.) We say that m and n are in the same generation when the distance from m to the root of PC is the same as the distance from n to the root of PC . Consider the following recursive query that computes the `sameGeneration` relation:

```
WITH RECURSIVE sameGeneration(m, n) AS
((SELECT parent, parent FROM PC)
UNION
(select child, child from PC)
UNION
SELECT t1.child, t2.child
FROM sameGeneration pair, PC t1, PC t2
WHERE pair.m = t1.parent and pair.n = t2.parent)
select distinct pair.m, pair.n from sameGeneration pair order by m, n;
```

Write a non-recursive function `sameGeneration()` in the language PL/pgSQL that computes the `sameGeneration` relation.

[25 points]